

Supplementary Material: Hamiltonian-Inspired Energy Dissipation and Chunk-Wise Variational State Coupling for Efficient Multimodal State-Space Models

Abhinav Sai Maddineni

I. OVERVIEW

This supplementary document provides additional architectural specifications, parameter counts, empirical diagnostic trajectories, dataset split manifests, and microsecond-level latency scaling measurements supporting the main manuscript.

II. COMPLETE PARAMETER BREAKDOWN

Table I presents the parameter allocation across all four experimental configurations. The frozen encoders (ViT-B/16 and RoBERTa-base) comprise 209.854 M parameters. The newly introduced trainable projection, recurrence, and coupling components require between 0.330 M and 0.422 M parameters ($< 0.21\%$ of total network capacity).

TABLE I
DETAILED PARAMETER BREAKDOWN ACROSS EXPERIMENTAL CONFIGURATIONS.

Configuration	Frozen (M)	Trainable (M)	Total (M)
SSD Baseline	209.854	0.330	210.184
w/o HEDO	209.854	0.356	210.209
w/o HVSC	209.854	0.397	210.251
Full HEDO-HVSC (Ours)	209.854	0.422	210.276

III. PRE-TRAINING DIAGNOSTIC AND COORDINATE TRAJECTORIES

Prior to full benchmark training, the numerical stability and coordinate-momentum behavior of HEDO were evaluated on initial batches.

A. Image Feature Transformation

- Raw Input Energy: $\mathcal{H}_{\text{raw}} = 170.824$
- Post-HEDO Energy: $\mathcal{H}_{\text{HEDO}} = 77.563$
- Cosine Similarity: $\text{sim}(\mathbf{X}_{\text{raw}}, \mathbf{X}_{\text{HEDO}}) = 0.9952$
- Unforced Trajectory over 5 steps: [170.824, 170.651, 170.927, 171.646, 172.801, 174.387]

B. Text Feature Transformation

- Raw Input Energy: $\mathcal{H}_{\text{raw}} = 5.849$
- Post-HEDO Energy: $\mathcal{H}_{\text{HEDO}} = 77.414$
- Cosine Similarity: $\text{sim}(\mathbf{X}_{\text{raw}}, \mathbf{X}_{\text{HEDO}}) = 0.9920$
- Unforced Trajectory over 5 steps: [5.849, 5.796, 5.774, 5.783, 5.823, 5.892]

C. Initial Layer Gradient Norms

- Custom SSD Recurrence: $\|\nabla_{\theta}\|_{\text{SSD}} = 2.3619$
- HEDO Operator: $\|\nabla_{\theta}\|_{\text{HEDO}} = 0.8861$
- HVSC Coupling Block: $\|\nabla_{\theta}\|_{\text{HVSC}} = 1.3847$

IV. DETAILED DATASET SPLIT MANIFEST

The official Flickr8k split was strictly enforced without data leakage:

- **Train Set:** 6,000 unique image filenames, 30,000 text descriptions.
- **Validation Set:** 1,000 unique image filenames, 5,000 text descriptions.
- **Held-Out Test Set:** 1,000 unique image filenames, 5,000 text descriptions.
- **Train** $\cap$ **Val** = $\emptyset$, **Train** $\cap$ **Test** = $\emptyset$, **Val** $\cap$ **Test** = $\emptyset$.

V. EMPIRICAL LATENCY SCALING DATA

Table II provides the exact microsecond-level timing measurements across sequence lengths $L \in \{16, 32, 64, 128, 256\}$ on an NVIDIA Tesla T4 GPU ($B = 16, d_{\text{model}} = 128, d_{\text{state}} = 64$).

TABLE II
EMPIRICAL LATENCY MEASUREMENTS FOR CUSTOM PYTORCH SSD RECURRENT SCAN.

Sequence Length (L)	Measured Latency (ms)	Standard Error (ms)	Linear M
16	2.89356	± 0.021	
32	5.43640	± 0.034	
64	10.31483	± 0.052	
128	19.71283	± 0.089	
256	38.60174	± 0.141	

Linear regression parameters:

$$\text{Latency}(L) = 0.149 \cdot L + 0.582 \text{ ms} \quad (R^2 = 0.9998) \quad (1)$$