



(a) Head Throughput vs. Retrieval Accuracy



(b) End-to-End Throughput vs. Peak Memory Footprint



Note: Bubble area represents trainable parameter count (465k to 764k parameters).